


Data import↓/export↑

Atomic vector (1d)



```
dept <- c("Physics", "Mathematics", "Statistics", "Computer Science")
nstaff <- c(12L, 8L, 20L, 23L)
```

image credit: Jenny Bryan

2 / 49

1d → 2d



```
library(tibble)
sci_tbl <- tibble(
  department = dept,
  count = nstaff,
  percentage = count / sum(count))
sci_tbl
```

```
#> # A tibble: 4 x 3
#>   department    count percentage
#>   <chr>         <int>      <dbl>
#> 1 Physics          12      0.190
#> 2 Mathematics       8      0.127
#> 3 Statistics       20      0.317
#> 4 Computer Science  23      0.365
```

image credit: Jenny Bryan

3 / 49

Beyond 1d vectors

1. Lists
2. Matrices and arrays
3. Data frames and tibbles

4 / 49


data strs
- lists

An object contains elements of **different data types**.

```
lst <- list( # list constructor/creator
  1:3,
  "a",
  c(TRUE, FALSE, TRUE),
  c(2.3, 5.9)
)
lst

#> [[1]]
#> [1] 1 2 3
#>
#> [[2]]
#> [1] "a"
#>
#> [[3]]
#> [1] TRUE FALSE TRUE
#>
#> [[4]]
#> [1] 2.3 5.9
```

data strs
- lists



data type data structure

```
typeof(lst) # primitive type
str(lst)
# el can be of diff lengths

#> [1] "list"

#> List of 4
#> $ : int [1:3] 1 2 3
#> $ : chr "a"
#> $ : logi [1:3] TRUE FALSE TRUE
#> $ : num [1:2] 2.3 5.9

class(lst) # type + attributes

#> [1] "list"
```

data strs
- lists

```
lst

#> [[1]]
#> [1] 1 2 3
#>
#> [[2]]
#> [1] "a"
#>
#> [[3]]
#> [1] TRUE FALSE TRUE
#>
#> [[4]]
#> [1] 2.3 5.9
```



data strs
- lists

A list can contain other lists, i.e. **recursive**

```
# a named list
str(list(first_el = lst, second_el = mtcars))

#> List of 2
#> $ first_el :list of 4
#> ..$ : int [1:3] 1 2 3
#> ..$ : chr "a"
#> ..$ : logi [1:3] TRUE FALSE TRUE
#> ..$ : num [1:2] 2.3 5.9
#> $ second_el :data.frame': 32 obs. of 11 variables:
#> ..$ mpg : num [1:32] 21 21 22.8 21.4 18.7 18.1 14.3 24.4 22.8 19
#> ..$ cyl : num [1:32] 6 6 4 6 8 6 8 4 4 6 ...
#> ..$ disp : num [1:32] 160 160 108 258 360 ...
#> ..$ hp : num [1:32] 110 110 93 110 175 105 245 62 95 123 ...
#> ..$ drat : num [1:32] 3.9 3.9 3.85 3.08 3.15 2.76 3.21 3.69 3.92 ...
#> ..$ wt : num [1:32] 2.62 2.88 2.32 3.21 3.44 ...
#> ..$ qsec : num [1:32] 16.5 17 18.6 19.4 17 ...
#> ..$ vs : num [1:32] 0 0 1 1 0 1 0 1 1 1 ...
#> ..$ am : num [1:32] 1 1 1 0 0 0 0 0 0 ...
#> ..$ gear : num [1:32] 4 4 4 3 3 3 3 4 4 4 ...
```

data str
- lists

Test for a list

```
is.list(lst)
```

```
#> [1] TRUE
```

Coerce to a list

```
as.list(1:3)
```

```
#> [[1]]  
#> [1] 1  
#>  
#> [[2]]  
#> [1] 2  
#>  
#> [[3]]  
#> [1] 3
```

data str
- lists
- matrices

2D structure of homogeneous data types

➤ `matrix()` to construct a matrix

```
matrix(1:9, nrow = 3)
```

```
#>      [,1] [,2] [,3]  
#> [1,]    1    4    7  
#> [2,]    2    5    8  
#> [3,]    3    6    9
```

➤ `as.matrix()` to coerce to a matrix

➤ `is.matrix()` to test for a matrix

data str
- lists

Subset by `[]`

```
lst[1]
```

```
#> [[1]]  
#> [1] 1 2 3
```

Subset by `[[]]`

```
lst[[1]]
```

```
#> [1] 1 2 3
```




image credit: Hadley Wickham

data str
- lists
- matrices

array: more than 2D matrix

```
array(1:9, dim = c(1, 3, 3))
```

```
#> , , 1  
#>  
#>      [,1] [,2] [,3]  
#> [1,]    1    2    3  
#>  
#> , , 2  
#>  
#>      [,1] [,2] [,3]  
#> [1,]    4    5    6  
#>  
#> , , 3  
#>  
#>      [,1] [,2] [,3]  
#> [1,]    7    8    9
```

data str

- lists
- matrices
- tibbles

A data frame is a **named list** of vectors of the **same length**.

```
sci_df <- data.frame(
  department = dept,
  count = nstaff)
sci_df

#>      department count
#> 1      Physics    12
#> 2    Mathematics     8
#> 3      Statistics    20
#> 4 Computer Science    23
```

data str

- lists
- matrices
- tibbles

The underlying data type is a list.

```
typeof(sci_df)

#> [1] "list"
```

data class	data attributes (meta info)
<pre>class(sci_df)</pre>	<pre>attributes(sci_df)</pre>
<pre>#> [1] "data.frame"</pre>	<pre>#> \$names #> [1] "department" "count" #> \$class #> [1] "data.frame" #> \$row.names #> [1] 1 2 3 4</pre>

data str

- lists
- matrices
- tibbles

A tibble is a **modern reimagining** of the data frame.


```
library(tibble)
sci_tbl <- tibble(
  department = dept,
  count = nstaff,
  percentage = count / sum(count))
sci_tbl

#> # A tibble: 4 x 3
#>   department    count percentage
#>   <chr>      <int>     <dbl>
#> 1 Physics         12     0.190
#> 2 Mathematics      8     0.127
#> 3 Statistics      20     0.317
#> 4 Computer Science 23     0.365
```

- as_tibble() to coerce to a tibble
- is_tibble() to test for a tibble

data str

- lists
- matrices
- tibbles



```
typeof(sci_tbl) # list in essence

#> [1] "list"
```

```
class(sci_tbl) # tibble is a special class of data.frame

#> [1] "tbl_df"      "tbl"        "data.frame"
```

Why tibble not data frame?

```
sci_df <- data.frame(
  department = dept,
  count = nstaff)
sci_df
```

```
#>      department count
#> 1      Physics     12
#> 2    Mathematics      8
#> 3      Statistics    20
#> 4 Computer Science    23
```

```
sci_tbl <- tibble(
  department = dept,
  count = nstaff,
  percentage = count / sum(count))
sci_tbl
```

```
#> # A tibble: 4 x 3
#>   department     count percentage
#>   <chr>         <int>         <dbl>
#> 1 Physics             12         0.190
#> 2 Mathematics          8         0.127
#> 3 Statistics          20         0.317
#> 4 Computer Science     23         0.365
```

Glimpse data

```
glimpse(sci_tbl) # to replace str()
```

```
#> Rows: 4
#> Columns: 3
#> $ department <chr> "Physics", "Mathematics", "Statistics",...
#> $ count         <int> 12, 8, 20, 23
#> $ percentage <dbl> 0.1904762, 0.1269841, 0.3174603, 0.3650...
```

Data types and their abbreviations

- chr: character
- dbl: double
- int: integer
- lgl: logical
- fct: factor
- date: date
- dtm: date-time
- more column data types

Subsetting tibble

- to 1d

with [[]] or \$

```
sci_tbl[["count"]] # col name
```

```
#> [1] 12 8 20 23
```

```
sci_tbl[[2]] # col pos
```

```
#> [1] 12 8 20 23
```

```
sci_tbl$count # col name
```

```
#> [1] 12 8 20 23
```

Subsetting tibble

- to 1d

- by columns

with [] or [, col]

```
sci_tbl["count"]
```

```
#> # A tibble: 4 x 1
#>   count
#>   <int>
#> 1    12
#> 2     8
#> 3    20
#> 4    23
```

```
sci_tbl[2] # sci_tbl[, 2]
```

```
#> # A tibble: 4 x 1
#>   count
#>   <int>
#> 1    12
#> 2     8
#> 3    20
#> 4    23
```

Subsetting tibble

- to 1d
- by columns
- by rows

```
sci_tbl[c(1, 3), ]
```

```
sci_tbl[-c(2, 4), ]
```

```
#> # A tibble: 2 x 3
#>   department count percentage
#>   <chr>      <int>      <dbl>
#> 1 Physics      12      0.190
#> 2 Statistics    20      0.317
```

```
#> # A tibble: 2 x 3
#>   department count percentage
#>   <chr>      <int>      <dbl>
#> 1 Physics      12      0.190
#> 2 Statistics    20      0.317
```

21 / 49

Subsetting tibble

- to 1d
- by columns
- by rows
- by cols & rows

```
sci_tbl[1:3, 2]
```

```
## sci_tbl[-4, 2]
## sci_tbl[1:3, "count"]
## sci_tbl[c(rep(TRUE, 3), FALSE), 2]
```

```
#> # A tibble: 3 x 1
#>   count
#>   <int>
#> 1    12
#> 2     8
#> 3    20
```

22 / 49

Subsetting tibble

- Use `[[` to extract 1d vectors from 2d tibbles
- Use `[` to subset tibbles to a new tibble
 - numbers (positive/negative) as indices
 - characters (column names) as indices
 - logicals as indices

```
sci_tbl[1:3, 2]
sci_tbl[-4, 2]
sci_tbl[1:3, "count"]
sci_tbl[c(rep(TRUE, 3), FALSE), 2]
```

23 / 49

The **tidyverse** is an opinionated collection of R packages designed for data science. *All packages share an underlying design philosophy, grammar, and data structures.*

24 / 49


Use {tidyverse}

```
library(tidyverse)


#> — Attaching packages — tidyverse 1.3.0 —
#> ✓ ggplot2 3.3.3   ✓ purrr  0.3.4
#> ✓ tibble  3.1.0   ✓ dplyr  1.0.5
#> ✓ tidyr   1.1.3   ✓ stringr 1.4.0
#> ✓ readr   1.4.0   ✓ forcats 0.5.1

#> — Conflicts — tidyverse_conflicts() —
#> * dplyr::filter() masks stats::filter()
#> * dplyr::lag()    masks stats::lag()
```

Data import



26 / 49



Reading plain-text rectangular files

(a.k.a. flat or spreadsheet-like files)

- > delimited text files with read_delim()
 - >> .csv: comma separated values with read_csv()
 - >> .tsv: tab separated values read_tsv()
- > .fwf: fixed width files with read_fwf()

```
head -4 data/pisa/pisa-student.csv # shell command, not R
```

```
#> year country school_id student_id mother_educ father_educ gender c
#> 2000 ALB 1001 1 NA NA female NA no 324.35 397.87 345.66 2.16 yes n
#> 2000 ALB 1001 3 NA NA female NA no NA 368.41 385.83 2.16 yes yes n
#> 2000 ALB 1001 6 NA NA male NA no NA 294.17 327.94 2.16 yes yes no
```



Reading comma delimited files

```
library(readr) # library(tidyverse)
pisa <- read_csv("data/pisa/pisa-student.csv", n_max = 2929621)
pisa
```

```
#> # A tibble: 2,929,621 x 22
#>   year country school_id student_id mother_educ father_educ
#>   <dbl> <chr>      <dbl>      <dbl> <dbl> <dbl>
#> 1  2000 ALB         1001          1 NA      NA
#> 2  2000 ALB         1001          3 NA      NA
#> 3  2000 ALB         1001          6 NA      NA
#> 4  2000 ALB         1001          8 NA      NA
#> 5  2000 ALB         1001         11 NA      NA
#> 6  2000 ALB         1001         12 NA      NA
#> # ... with 2,929,615 more rows, and 16 more variables:
#> #   gender <chr>, computer <dbl>, internet <chr>,
#> #   math <dbl>, read <dbl>, science <dbl>, stu_wgt <dbl>,
#> #   desk <chr>, room <chr>, dishwasher <chr>,
#> #   television <chr>, computer_n <chr>, car <chr>,
#> #   book <chr>, wealth <dbl>, escs <dbl>
```

29 / 49

Let's talk about the file path again!

```
pisa <- read_csv("data/pisa/pisa-student.csv", n_max = 2929621)
```

data/pisa/pisa-student.csv relative to the top-level (or root) directory:

```
> stats220.Rproj
> data/
>> pisa/pisa-student.csv
```

If you don't like `/`, you can use here: `here()` instead.

```
read_csv(here::here("data", "pisa", "pisa-student.csv"))
```

NOTE: I use the `here()` function from the `(here)` package using `pkg::fun()`, without calling `library(here)` the usual way. 49




read_csv() arguments with ?read_csv()

```
read_csv(
  file,
  col_names = TRUE,
  col_types = NULL,
  locale = default_locale(),
  na = c("", "NA"),
  quoted_na = TRUE,
  quote = "\"",
  comment = "#",
  trim_ws = TRUE,
  skip = 0,
  n_max = Inf,
  guess_max = min(1000, n_max),
  progress = show_progress(),
  skip_empty_rows = TRUE
)
```

31 / 49



Faster delimited reader at 1.4GB/sec



```
library(vroom)
pisa <- vroom("data/pisa/pisa-student.csv", n_max = 2929621)
```

32 / 49



Reading proprietary binary files

- Microsoft Excel
 - .xls: MSFT Excel 2003 and earlier
 - .xlsx: MSFT Excel 2007 and later

```
library(readxl)
time_use <- read_xlsx("data/time-use-oecd.xlsx")
time_use
```

```
#> # A tibble: 461 x 3
#>   Country Category `Time (minutes)`
#>   <chr>    <chr>          <dbl>
#> 1 Australia Paid work         211.
#> 2 Austria  Paid work         280.
#> 3 Belgium  Paid work         194.
#> 4 Canada   Paid work         269.
#> 5 Denmark  Paid work         200.
#> 6 Estonia  Paid work         231.
#> # _ with 455 more rows
```

33 / 49



Reading proprietary binary files

- SAS
 - .sas7bdat with read_sas()
- Stata
 - .dta with read_dta()
- SPSS
 - .sav with read_sav()

```
library(haven)
pisa2018 <- read_spss("data/pisa/CY07_MSU_STU_QQQ.sav")
```

Raw PISA data is made available in SAS and SPSS data formats.

data source: <https://www.oecd.org/pisa/data/2018database/>

34 / 49

Your turn

“

What is the R data format for a single object? What is its file extension?

35 / 49

Well, SQL!

- **Structured Query Language** for accessing and manipulating databases.
- Relational database management systems
 - SQLite
 - MySQL
 - PostgreSQL
 - BigQuery
 - Spark SQL

However, 220 is all about R!

36 / 49

{DBI}

Connecting R to database*

```
library(RSQLite)
con <- dbConnect(SQLite(), dbname = "data/pisa/pisa-student.db")
dbListTables(con)

#> [1] "pisa"

dbListFields(con, "pisa")

#> [1] "year"      "country"    "school_id"  "student_id" "moth
#> [6] "father_educ" "gender"     "computer"   "internet"   "math
#> [11] "read"       "science"    "stu_wgt"    "desk"       "room
#> [16] "dishwasher" "television" "computer_n" "car"        "book
#> [21] "wealth"     "escs"
```

NOTE: slides marked with * are not examinable.

37 / 49

{DBI}

Connecting R to database*

➤ reading data from database

```
pisa <- dbReadTable(con, "pisa")
```

➤ writing SQL queries to read chunks

```
res <- dbSendQuery(con, "SELECT * FROM pisa WHERE year = 2018")
pisa2018 <- dbFetch(res)
```

➤ closing connection

```
dbDisconnect(con)
```

38 / 49



Reading chunks for larger than memory data*

```
chunked <- function(x, pos) {
  dplyr::filter(x, year == 2018)
}
pisa2018 <- read_csv_chunked("data/pisa/pisa-student.csv",
  callback = DataFrameCallback$new(chunked))
```

39 / 49

{jsonlite}

JSON: JavaScript Object Notation

- object: {}
- array: []
- value: string/character, number, object, array, logical, null

JSON

```
{
  "firstName": "Eero",
  "lastName": "Wang",
  "address": {
    "city": "Auckland",
    "postalCode": 1010
  },
  "logical": [true, false]
}
```

R list

```
list(
  firstName = "Eero",
  lastName = "Wang",
  address = list(
    city = "Auckland",
    postalCode = 1010
  ),
  logical = c(TRUE, FALSE)
)
```

40 / 49

{jsonlite}

Reading json files

```
library(jsonlite)
url <- "https://vega.github.io/vega-editor/app/data/movies.json"
movies <- read_json(url)
length(movies)

#> [1] 3201

movies[[1]]

#> $Title
#> [1] "The Land Girls"
#>
#> $US_Gross
#> [1] 146083
#>
#> $Worldwide_Gross
#> [1] 146083
#>
#> $US_DVD_Sales
#> [1] 146083
```

{jsonlite}

Reading json files as tibbles

```
movies_tbl <- as_tibble(read_json(url, simplifyVector = TRUE))
movies_tbl

#> # A tibble: 3,201 x 16
#>   Title US_Gross Worldwide_Gross US_DVD_Sales
#>   <chr>      <int>      <dbl>      <int>
#> 1 The Land Girls 146083 146083 NA
#> 2 First Love, Last Rites 10876 10876 NA
#> 3 I Married a Strange Person 203134 203134 NA
#> 4 Let's Talk About Sex 373615 373615 NA
#> 5 Slam 1009819 1009819 NA
#> 6 Mississippi Mermaid 24551 2624551 NA
#> # _ with 3,195 more rows, and 12 more variables:
#> #   Production_Budget <int>, Release_Date <chr>,
#> #   MPAA_Rating <chr>, Running_Time_min <int>,
#> #   Distributor <chr>, Source <chr>, Major_Genre <chr>,
#> #   Creative_Type <chr>, Director <chr>,
#> #   Rotten_Tomatoes_Rating <int>, IMDB_Rating <dbl>,
#> #   IMDB_Votes <int>
```



Reading spatial data*

```
library(sf)
aki_bus <- st_read("data/BusService/BusService.shp")

#> Reading layer 'BusService' from data source `~/Users/wany568/Teaching/
#> Simple feature collection with 509 features and 7 fields
#> geometry type: MULTILINESTRING
#> dimension: XY
#> bbox: xmin: 1727652 ymin: 5859539 xmax: 1787138 ymax: 5910000
#> projected CRS: NZGD2000_New_Zealand_Transverse_Mercator_2000
```



Reading spatial data*

```
library(sf)
aki_bus <- st_read("data/BusService/BusService.shp")

#> Reading layer 'BusService' from data source `~/Users/wany568/Teaching/
#> Simple feature collection with 509 features and 7 fields
#> geometry type: MULTILINESTRING
#> dimension: XY
#> bbox: xmin: 1727652 ymin: 5859539 xmax: 1787138 ymax: 5910000
#> projected CRS: NZGD2000_New_Zealand_Transverse_Mercator_2000
```



Reading spatial data*

```
akl_bus[1:4, ]
```


```
#> Simple feature collection with 4 features and 7 fields
#> geometry type:  MULTILINESTRING
#> dimension:      XY
#> bbox:           xmin: 1751253 ymin: 5915245 xmax: 1758019 ymax: 5915245
#> projected CRS:  NZGD2000_New_Zealand_Transverse_Mercator_2000
#>   OBJECTID ROUTEPATID AGENCYNAME ROUTENUMBE MODE Shape__Len geometry
#> 1    343077      02805      NZB St Lukes To Wynyard Quarter Via K 20      7948.418 MULTILINESTRING ((1755487 5...
#> 2    343078      02806      NZB Wynyard Quarter To St Lukes Via K 20      7919.198 MULTILINESTRING ((1756321 5...
#> 3    343079      02209      NZB Avondale To City Centre Via New 22A     11419.588 MULTILINESTRING ((1757613 5...
#> 4    343080      02208      NZB City Centre To Avondale Via New 22A     11607.711 MULTILINESTRING ((1757346 5...
```

45 / 49



Spatial visualisation*

[Map](#) [R Code](#)



46 / 49

Data export

47 / 49

From read_*() to write_*()

```
write_csv(movies_tbl, file = "data/movies.csv")
write_sas(movies_tbl, path = "data/movies.sas7bdat")
write_json(movies_tbl, path = "data/movies.json")
```

48 / 49

Reading



- Tibbles
- Data import



- Subsetting